


Practical Concurrent and Parallel Programming XII

Message Passing I

Raúl Pardo

Please participate in the course evaluation





- Problems in shared memory concurrency (revisited)
- Actors
- Erlang
- Example systems
 - Turnstile (counter)
 - Broadcaster
 - Bounded Buffer

Problems in shared memory concurrency



“Writing thread-safe code is, at its core, about managing access to shared mutable data”

Goetz

Problems in shared memory concurrency



“Writing thread-safe code is, at its core, about managing access to shared mutable data”

Goetz

What problems have we seen in concurrent access to shared memory?



“Writing thread-safe code is, at its core, about managing access to shared mutable data”

Goetz

- Race conditions
- Data races
- Visibility
- Reasoning is tricky
 - Specially lock-free computation 😊

Problems in shared memory concurrency



“Writing thread-safe code is, at its core, about managing access to shared mutable data”

Goetz

What solutions have we seen to the problems in concurrent access to shared memory?



“Writing thread-safe code is, at its core, about managing access to shared mutable data”

Goetz

- Happens-before reasoning
- Linearizability reasoning
- For race conditions and data races:
 - Ensuring mutual exclusion
 - Trying to ensure progress
 - Immutability
 - Compare and Swap (CAS) algorithms
 - Trying to ensure progress
- For visibility:
 - Volatile and final variables, idioms for safe publication, etc



“Writing thread-safe code is, at its core, about managing access to shared mutable data”

Goetz

- Happens-before reasoning
- Linearizability reasoning
- For race conditions and data races:
 - Ensuring mutual exclusion
 - Trying to ensure progress
 - Immutability
 - Compare and Swap (CAS) algorithms
 - Trying to ensure progress
- For visibility:
 - Volatile and final variables, idioms for safe publication, etc.

Why don't we simply avoid sharing state?
This is the idea behind message passing!

Message passing concurrency



- Threads do not share state
- If threads need to share data, then data must be communicated by sending messages
- Threads work only on their own local memory

One of the
designers of Erlang



Joe Armstrong
@joeerl

Following



Copying = good, sharing=bad



Hey @joeerl, do you think the inter-process communication should never be done by sharing memory? Otherwise, when it's okay?
Thanks a lot!

12:11 PM - 22 Nov 2018

11 Retweets 18 Likes




Shared memory vs Message Passing




- Shared Memory

- Synchronisation by writing in shared memory



- Message Passing

- Synchronisation by sending messages

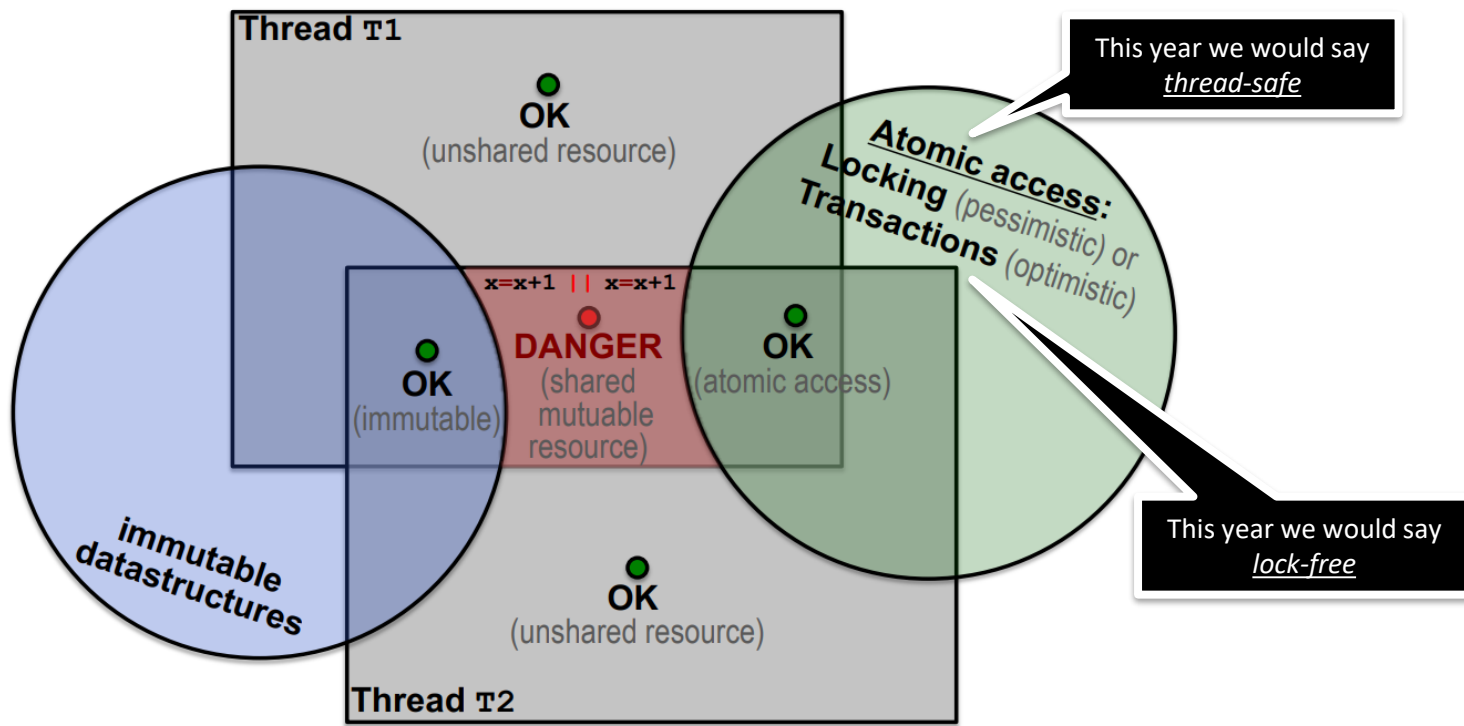




PROBLEM: **Sharing && Mutability!**

SOLUTIONS:

- 1) atomic access!
locking or transactions
NB: avoid deadlock!
- 2) avoid mutability!
- 3) avoid sharing...

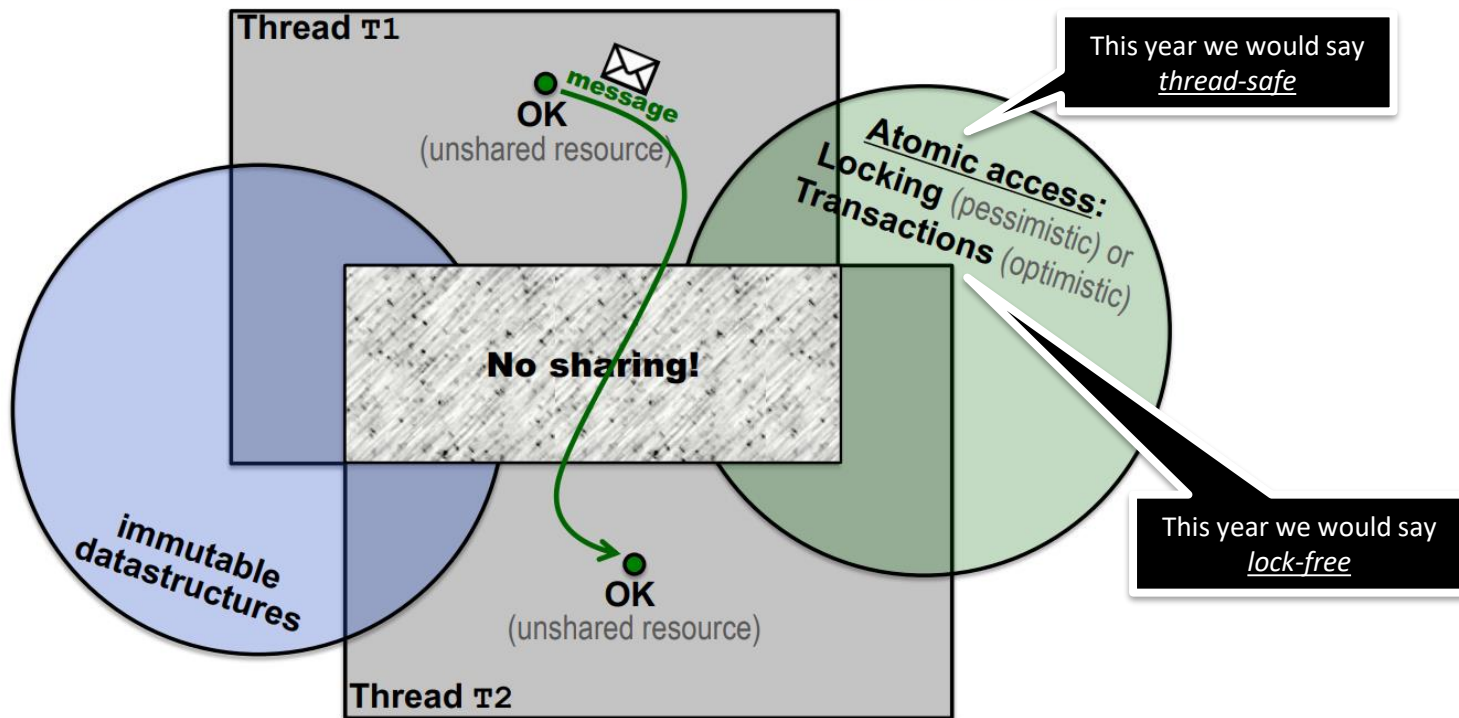




PROBLEM: **Sharing && Mutability!**

SOLUTIONS:

- 1) atomic access!
locking or transactions
NB: avoid deadlock!
- 2) avoid mutability!
- 3) avoid sharing...




- How should we implement message passing concurrency?
- A possible solution is use standard communication systems
 - Sockets (last week)
 - Remote Procedure Calls (RPC)
 - Java Remote Method Invocation (RMI)
 - Message passing interfaces (MPI)

combined with concurrency as we have seen so far

Message passing concurrency



- How should we implement message passing concurrency?
- Another option is to *use a concurrency model with message passing built-in*
 - That is, the *actor model*!
- The actors model was first introduced by [Hewitt'73] and later formalized by [Agha'85] (part of the readings)
 - [Hewitt'73] - Carl Hewitt, Peter Bishop & Richard Steiger. A universal modular ACTOR formalism for artificial intelligence. IJCAI'73: Proceedings of the 3rd international joint conference on Artificial intelligence. 1973.
 - [Agha'85] – Gul A. Agha. ACTORS: A Model of Concurrent Computation in Distributed Systems. MIT Press. 1985.




Actor model

What is an Actor? (Bird's eye)




- An actor can be seen as a sequential unit of computation
 - Although, formally, the model allows for parallelism within the actor, one can safely assume that there are not concurrency issues within the actor.
 - You can think of an actor as a thread
- Actors can send messages to other actors






- An actor is an abstraction of a thread (intuitively)
- An actors can only execute any of these 4 actions
 1. Receive messages from other actors
 2. Send asynchronous messages to other actors
 3. Create new actors
 4. Change its behaviour (local state and/or message handlers)
- Actors do not share memory
 - They only have access to:
 - Their *local state* (local memory)
 - Their *mailbox* (multiset with received messages)
 - By default: i) the mailbox is unbounded, but with a fixed size at each point of the execution;
ii) messages are ordered in a FIFO style



- Every actor in the system has a unique identifier
 - A.k.a. mail address or actor reference
- Actors can
 - Send (finitely many) messages
 - Receive (finitely many) message
 - Received messages are placed in the actor's mailbox (asynchronous communication, see next slide)
- Messages include
 - Content of the message (arbitrary payload)

Asynchronous communication

· 19




- Asynchronous send:
 - The sender places the message in the mailbox of the receiver
 - It is non-blocking
- Asynchronous receive:
 - The receiver takes a message from the mailbox
 - The receiver blocks if the mailbox is empty

No requirements on message arrival order



- No assumptions should be made about the order of arrival of messages
- For instance, consider this sequence of operations
 1. Actor1 sends message M1 to Actor2
 2. Actor1 sends message M2 to Actor2
- It is not guaranteed that M1 arrives before M2



Erlang (Actors implementation)



- Developed by Joe Armstrong, Robert Virding, and Mike Williams in 1986
 - Open-sourced in 1998; despite Ericsson's attempts to prevent it
- Erlang = **E**ricsson **L**anguage
 - (Presumably) named after the Danish mathematician Agner Krarup Erlang (1878 –1929) for his pioneering and influential work in the field of telecommunications
- Language developed for telephony applications
 - Erlang/OTP is supported and maintained by the Open Telecom Platform (OTP) product unit at Ericsson.
- Famously used at WhatsApp (among many other companies)
 - In 2014, there were only 32 engineers at WhatsApp developing/maintaining software for 450 million users
- Multiple companies use Erlang in production
 - <https://erlang-companies.org/>

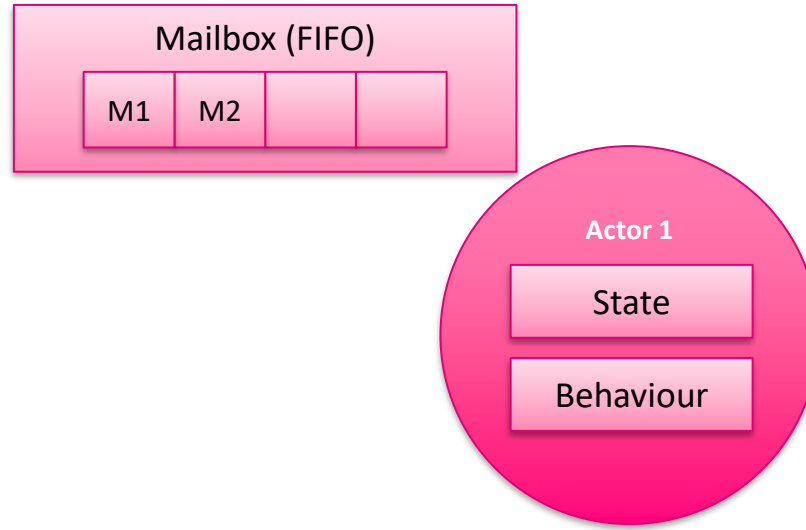


...



MOTOROLA SOLUTIONS

Erlang actors (processes)



Tivoli Turnstiles with Actors!

· 24




- Actors system to count the numbers of visitors in Tivoli



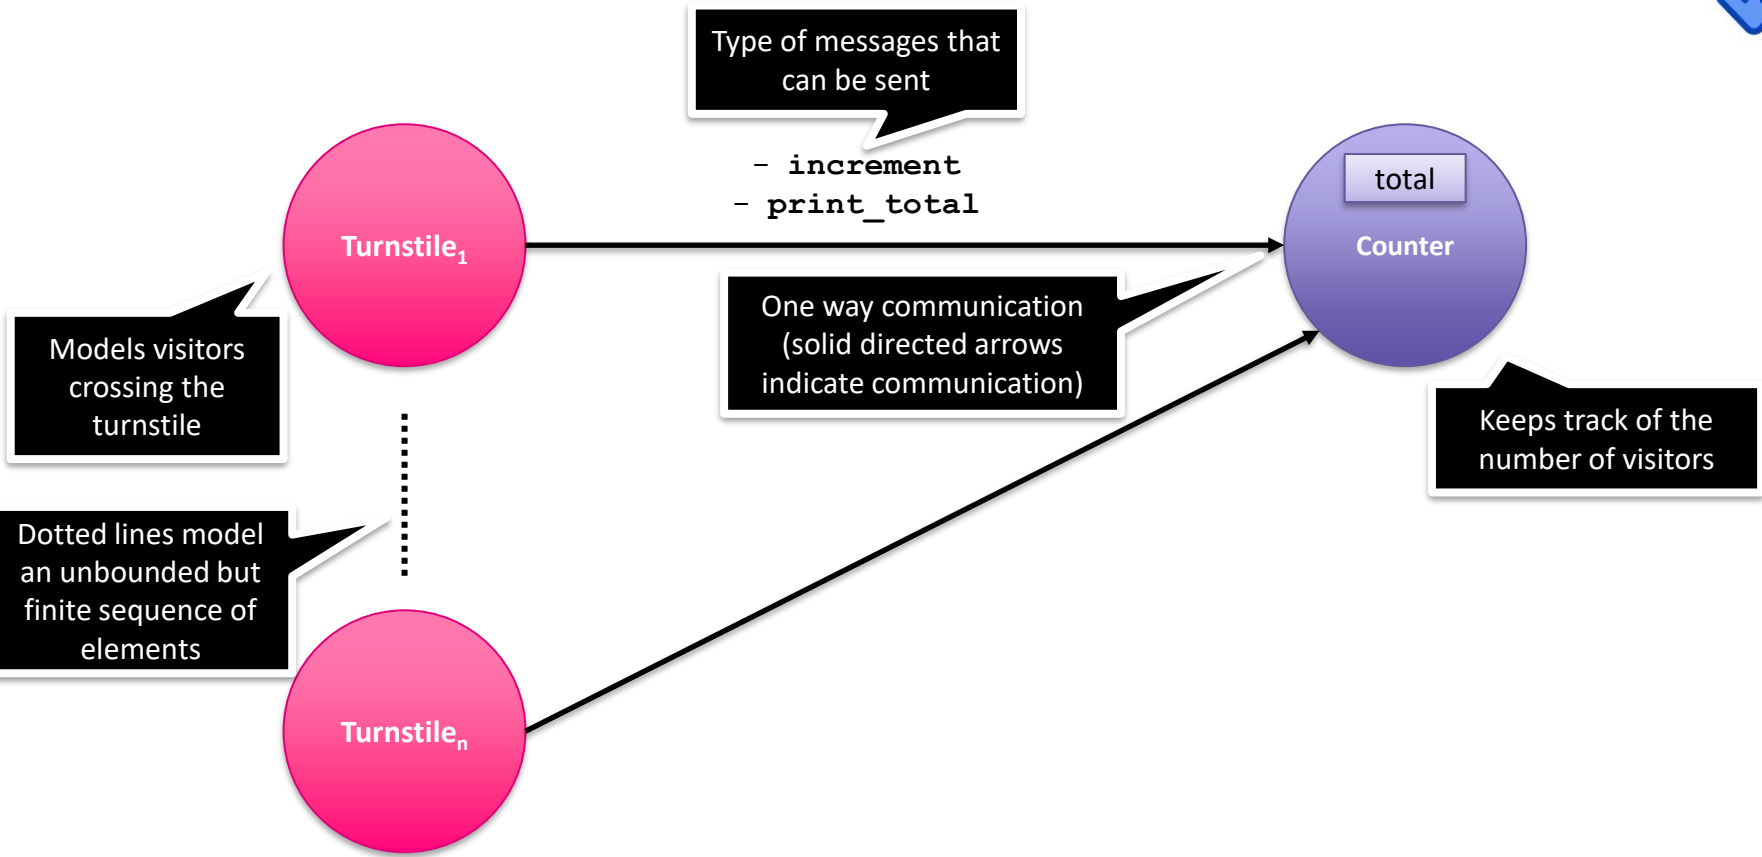
Turnstile with Actors - Design



· 25



Turnstile with Actors - Design



Turnstile with Actors - Implementation

· 26



```
-module(counter).  
-export([start/0, init/1, loop/1]).
```


```
% State of the actor  
-record(counter_state, {total}).
```

```
% Function to create counter actors  
start() ->  
    spawn(?MODULE, init, [0]).
```

```
% Function to initialize the state and the actors behavior  
% upon receiving messages  
init(InitialValue) ->  
    InitialState = #counter_state{total = InitialValue},  
    loop(InitialState).
```

```
% Function defining the behavior upon receiving messages  
loop(State) ->  
    receive  
        increment ->  
            CurrentTotal = State#counter_state.total,  
            NewState = State#counter_state{total = CurrentTotal + 1},  
            io:format("A visitor arrived 🍌~n"),  
            loop(NewState);  
        print_total ->  
            io:format("The counter value is ~p~n",  
                [State#counter_state.total]),  
            loop(State)  
    end.
```

- increment
- print_total



Turnstile with Actors - Implementation

· 27



```
-module(counter).  
-export([start/0, init/1, loop/1]).
```


```
% State of the actor  
-record(counter_state, {total}).
```

```
% Function to create counter actors  
start() ->  
    spawn(?MODULE, init, [0]).
```

```
% Function to initialize the state and the actors behavior  
% upon receiving messages  
init(InitialValue) ->  
    InitialState = #counter_state{total = InitialValue},  
    loop(InitialState).
```

```
% Function defining the behavior upon receiving messages  
loop(State) ->  
    receive  
        increment ->  
            CurrentTotal = State#counter_state.total,  
            NewState = State#counter_state{total = CurrentTotal + 1},  
            io:format("A visitor arrived 🍌~n"),  
            loop(NewState);  
        print_total ->  
            io:format("The counter value is ~p~n",  
                [State#counter_state.total]),  
            loop(State)  
    end.
```

- increment
- print_total



Actors vs Threads

Each actor is defined in its own module. This is similar to a Thread class for a Java thread.

Remember that the module name must match the file name, in this case `counter.erl`


Turnstile with Actors - Implementation

· 28



```
-module(counter).  
-export([start/0, init/1]).  
  
% State of the actor  
-record(counter_state, {total}).  
  
% Function to create counter actors  
start() ->  
    spawn(?MODULE, init, [0]).  
  
% Function to initialize the state and the actors behavior  
% upon receiving messages  
init(InitialValue) ->  
    InitialState = #counter_state{total = InitialValue},  
    loop(InitialState).  
  
% Function defining the behavior upon receiving messages  
loop(State) ->  
    receive  
        increment ->  
            CurrentTotal = State#counter_state.total,  
            NewState = State#counter_state{total = CurrentTotal + 1},  
            io:format("A visitor arrived 🍌~n"),  
            loop(NewState);  
        print_total ->  
            io:format("The counter value is ~p~n",  
                [State#counter_state.total]),  
            loop(State)  
    end.
```

- increment
- print_total



We must export the functions necessary to create and interact with the actor.

Turnstile with Actors - Implementation

· 29



```
-module(counter).  
-export([start/0, init/1, loop/1]).
```

```
% State of the actor  
-record(counter_state, {total}).
```

```
% Function to create counter actors  
start() ->  
    spawn(?MODULE, init, [0]).
```


```
% Function to initialize the state and the actors behavior  
% upon receiving messages  
init(InitialValue) ->  
    InitialState = #counter_state{total = InitialValue},  
    loop(InitialState).
```

```
% Function defining the behavior upon receiving messages  
loop(State) ->  
    receive  
        increment ->  
            CurrentTotal = State#counter_state.total,  
            NewState = State#counter_state{total = CurrentTotal + 1},  
            io:format("A visitor arrived 🍌~n"),  
            loop(NewState);  
        print_total ->  
            io:format("The counter value is ~p~n",  
                [State#counter_state.total]),  
            loop(State)  
    end.
```

Actors vs Threads

- Like threads actors' have a local state. The local state is commonly defined as a record.

- increment
- print_total



Turnstile with Actors - Implementation

· 29



```
-module(counter).  
-export([start/0, init/1, loop/1]).
```

```
% State of the actor  
-record(counter_state, {total}).
```

```
% Function to create counter actors  
start() ->  
    spawn(?MODULE, init, [0]).
```


```
% Function to initialize the state and the actors behavior  
% upon receiving messages  
init(InitialValue) ->  
    InitialState = #counter_state{total = InitialValue},  
    loop(InitialState).
```

```
% Function defining the behavior upon receiving messages  
loop(State) ->  
    receive  
        increment ->  
            CurrentTotal = State#counter_state.total,  
            NewState = State#counter_state{total = CurrentTotal + 1},  
            io:format("A visitor arrived 🍌~n"),  
            loop(NewState);  
        print_total ->  
            io:format("The counter value is ~p~n",  
                [State#counter_state.total]),  
            loop(State)  
    end.
```

Actors vs Threads

- Like threads actors' have a local state. The local state is commonly defined as a record.

- increment
- print_total



Are there visibility issues
in the actor state?

Turnstile with Actors - Implementation

· 30



```
-module(counter).  
-export([start/0, init/1, loop/1]).
```


```
% State of the actor  
-record(counter_state, {total}).
```

```
% Function to create counter actors  
start() ->  
    spawn(?MODULE, init, [0]).
```

```
% Function to initialize the state and the actors behavior  
% upon receiving messages  
init(InitialValue) ->  
    InitialState = #counter_state{total = InitialValue},  
    loop(InitialState).
```

```
% Function defining the behavior upon receiving messages  
loop(State) ->  
    receive  
        increment ->  
            CurrentTotal = State#counter_state.total,  
            NewState = State#counter_state{total = CurrentTotal + 1},  
            io:format("A visitor arrived 🍌~n"),  
            loop(NewState);  
        print_total ->  
            io:format("The counter value is ~p~n",  
                [State#counter_state.total]),  
            loop(State)  
    end.
```

- increment
- print_total



Actors vs Threads

- It is common to define a function `init` to initialize the state of the actor and its behavior (specified in the `loop` function that we will explain in the coming slides)

This is similar to the constructor of a Thread class for a Java thread.

Turnstile with Actors - Implementation

· 31



```
-module(counter).  
-export([start/0, init/1, loop/1]).
```

```
% State of the actor  
-record(counter_state, {total}).
```

```
% Function to create counter actors
```

```
start() ->  
    spawn(?MODULE, init, [0]).
```

```
% Function to initialize the state and the actors behavior
```


```
% upon receiving messages
```

```
init(InitialValue) ->  
    InitialState = #counter_state{total = InitialValue},  
    loop(InitialState).
```

```
% Function defining the behavior upon receiving messages
```

```
loop(State) ->  
    receive  
    increment ->  
        CurrentTotal = State#counter_state.total,  
        NewState = State#counter_state{total = CurrentTotal + 1},  
        io:format("A visitor arrived 🍌~n"),  
        loop(NewState);  
    print_total ->  
        io:format("The counter value is ~p~n",  
            [State#counter_state.total]),  
        loop(State)  
end.
```

- increment
- print_total



Actors vs Threads

- To start an actor, we use the built-in function (BIF) `spawn`. This function takes:

- an atom indicating the module where the `init` function is
- The `init` function
- a list of parameter for the `init` function

This is similar to the function `start()` in a Java Thread class.

Turnstile with Actors - Implementation

· 31



```
-module(counter).  
-export([start/0, init/1, loop/1]).
```

```
% State of the actor  
-record(counter_state, {total}).
```


```
% Function to create counter  
start() ->  
    spawn(?MODULE, init, [0]).
```

?MODULE is a predefined macro that expands to the module atom at compile time

```
% Function to initialize the state and the actors behavior  
% upon receiving messages  
init(InitialValue) ->  
    InitialState = #counter_state{total = InitialValue},  
    loop(InitialState).
```

```
% Function defining the behavior upon receiving messages  
loop(State) ->  
    receive  
    increment ->  
        CurrentTotal = State#counter_state.total,  
        NewState = State#counter_state{total = CurrentTotal + 1},  
        io:format("A visitor arrived 🍌~n"),  
        loop(NewState);  
    print_total ->  
        io:format("The counter value is ~p~n",  
            [State#counter_state.total]),  
        loop(State)  
    end.
```

- increment
- print_total



Actors vs Threads

- To start an actor, we use the built-in function (BIF) `spawn`. This function takes:

- an atom indicating the module where the `init` function is
- The `init` function
- a list of parameter for the `init` function

This is similar to the function `start()` in a Java Thread class.

Turnstile with Actors - Implementation

· 32



```
-module(counter).  
-export([start/0, init/1, loop/1]).
```


```
% State of the actor  
-record(counter_state, {total}).
```

```
% Function to create counter actors  
start() ->  
    spawn(?MODULE, init, [0]).
```

```
% Function to initialize the state and the actors behavior  
% upon receiving messages  
init(InitialValue) ->  
    InitialState = #counter_state{total = InitialValue},  
    loop(InitialState).
```

```
% Function defining the behavior upon receiving messages  
loop(State) ->  
    receive  
        increment ->  
            CurrentTotal = State#counter_state.total,  
            NewState = State#counter_state{total = CurrentTotal + 1},  
            io:format("A visitor arrived 🍌~n"),  
            loop(NewState);  
        print_total ->  
            io:format("The counter value is ~p~n",  
                [State#counter_state.total]),  
            loop(State)  
    end.
```

- increment
- print_total



- The loop function defines the behavior of the thread upon receiving messages
- It takes as a parameter the current state of the thread; as it might be needed in processing incoming messages.

Turnstile with Actors - Implementation

· 33



```
-module(counter).  
-export([start/0, init/1, loop/1]).
```

```
% State of the actor  
-record(counter_state, {total}).
```


```
% Function to create counter actors  
start() ->  
    spawn(?MODULE, init, [0]).
```

```
% Function to initialize the state and the actors behavior  
% upon receiving messages  
init(InitialValue) ->  
    InitialState = #counter_state{total = InitialValue},  
    loop(InitialState).
```

```
% Function defining the behavior upon receiving messages  
loop(State) ->  
    receive  
    increment ->  
        CurrentTotal = State#counter_state.total,  
        NewState = State#counter_state{total = CurrentTotal + 1},  
        io:format("A visitor arrived 🍌~n"),  
        loop(NewState);  
    print_total ->  
        io:format("The counter value is ~p~n",  
            [State#counter_state.total]),  
        loop(State)  
    end.
```

- The body of the loop function is a receive statement with one case per type of message that the actor can handle.

- increment
- print_total



- **receive** is a blocking statement that waits for incoming messages to the actor's mailbox that match any of the specified cases.

- Messages that do not match the cases in receive may add up in the mailbox, which may lead to slow processing times.

Turnstile with Actors - Implementation

· 34



```
-module(counter).  
-export([start/0, init/1, loop/1]).
```


```
% State of the actor  
-record(counter_state, {total}).
```

```
% Function to create counter actors  
start() ->  
    spawn(?MODULE, init, [0]).
```

```
% Function to initialize the state and the actors behavior  
% upon receiving messages  
init(InitialValue) ->  
    InitialState = #counter_state{total = InitialValue},  
    loop(InitialState).
```

```
% Function defining the behavior upon receiving messages  
loop(State) ->  
    receive  
    increment ->  
        CurrentTotal = State#counter_state.total,  
        NewState = State#counter_state{total = CurrentTotal + 1},  
        io:format("A visitor arrived 🍷~n"),  
        loop(NewState);  
    print_total ->  
        io:format("The counter value is ~p~n",  
            [State#counter_state.total]),  
        loop(State)  
    end.
```

- increment
- print_total



- For **increment** messages we increment the value of the counter
- For the **print_total** message we print the value

Turnstile with Actors - Implementation

· 35



```
-module(counter).  
-export([start/0, init/1, loop/1]).
```


```
% State of the actor  
-record(counter_state, {total}).
```

```
% Function to create counter actors  
start() ->  
    spawn(?MODULE, init, [0]).
```

```
% Function to initialize the state and the actors behavior  
% upon receiving messages  
init(InitialValue) ->  
    InitialState = #counter_state{total = InitialValue},  
    loop(InitialState).
```

```
% Function defining the behavior upon receiving messages  
loop(State) ->  
    receive  
        increment ->  
            CurrentTotal = State#counter_state.total,  
            NewState = State#counter_state{total = CurrentTotal + 1},  
            io:format("A visitor arrived 🍌~n"),  
            loop(NewState);  
        print_total ->  
            io:format("The counter value is ~p~n",  
                [State#counter_state.total]),  
            loop(State)  
    end.
```

- increment
- print_total



- For **increment**, we wait again for messages but using the **NewState**
- For the **print_total**, we wait for with the old state **State**

Turnstile with Actors - Implementation

· 36




```
-module(turnstile).  
-export([start/1, init/1, loop/1]).
```

```
% State of the actor  
-record(turnstile_state, {counter_server}).
```

```
% Function to create turnstile actors  
start(CounterPID) ->  
    spawn(?MODULE, init, [CounterPID]).
```

```
% Function to initialize the state and the actors behavior  
% upon receiving messages  
init(CounterPID) ->  
    State = #turnstile_state{counter_server = CounterPID},  
    loop(State).
```

```
% Function defining the behavior upon receiving messages  
loop(State) ->  
    receive  
        person_crossing ->  
            State#turnstile_state.counter_server ! increment,  
            loop(State)  
    end.
```



- The state of the turnstile actor is simply a variable to store the identifier of the actor, typically referred to in Erlang as “process identifier” (PID)

Turnstile with Actors - Implementation

· 37




```
-module(turnstile).  
-export([start/1, init/1, loop/1]).
```

```
% State of the actor  
-record(turnstile_state, {counter_server}).
```

```
% Function to create turnstile actors  
start(CounterPID) ->  
    spawn(?MODULE, init, [CounterPID]).
```

```
% Function to initialize the state and the actors behavior  
% upon receiving messages  
init(CounterPID) ->  
    State = #turnstile_state{counter_server = CounterPID},  
    loop(State).
```

```
% Function defining the behavior upon receiving messages  
loop(State) ->  
    receive  
        person_crossing ->  
            State#turnstile_state.counter_server ! increment,  
            loop(State)  
    end.
```



- Initialization is the same as in the counter actor. The only difference is that the start and init functions now take an initialization parameter from the user, namely the CounterPID, which refers to the PID of the counter process.

Turnstile with Actors - Implementation

· 38




```
-module(turnstile).  
-export([start/1, init/1, loop/1]).
```

```
% State of the actor  
-record(turnstile_state, {counter_server}).
```

```
% Function to create turnstile actors  
start(CounterPID) ->  
    spawn(?MODULE, init, [CounterPID]).
```

```
% Function to initialize the state and the actors behavior  
% upon receiving messages  
init(CounterPID) ->  
    State = #turnstile_state{counter_server = CounterPID},  
    loop(State).
```

```
% Function defining the behavior upon receiving messages  
loop(State) ->  
    receive  
        person_crossing ->  
            State#turnstile_state.counter_server ! increment,  
            loop(State)  
    end.
```



- In the loop function, we simply consider one case modelling a person crossing the turnstile.
- In this case, the turnstile actor sends an increment message to the counter server.
- This is done using the operation **PID ! Message** that sends an asynchronous message, **Message** to the process with identifier **PID**

Turnstile with Actors - Implementation

· 38



```
-module(turnstile).  
-export([start/1, init/1, loop/1]).
```


```
% State of the actor  
-record(turnstile_state, {counter_server}).
```

```
% Function to create turnstile actors  
start(CounterPID) ->  
    spawn(?MODULE, init, [CounterPID]).
```

```
% Function to initialize the state and the actors behavior  
% upon receiving messages  
init(CounterPID) ->  
    State = #turnstile_state{counter_server = CounterPID},  
    loop(State).
```

```
% Function defining the behavior upon receiving messages  
loop(State) ->  
    receive  
        person_crossing ->  
            State#turnstile_state.counter_server ! increment,  
            loop(State)  
    end.
```

Let's run the code in the `turnstile` folder!



- In the loop function, we simply consider one case modelling a person crossing the turnstile.
- In this case, the turnstile actor sends an increment message to the counter server.
- This is done using the operation `PID ! Message` that sends an asynchronous message, `Message` to the process with identifier `PID`

A technical note on receive in Erlang

· 39



- Messages when executing receive match from top to bottom
 - Similar to function cases
- Furthermore, the mailbox always preserves the order in which messages arrive, i.e., FIFO

```
loop() ->  
  receive  
    dog ->    ...;  
    cat ->    ...;  
    snake ->  ...;  
    Animal -> ...  
  end,  
  loop().
```

Mailbox (FIFO)

cat	dog	owl	
-----	-----	-----	--


A technical note on receive in Erlang

· 40



- Messages when executing receive match from top to bottom
 - Similar to function cases
- Furthermore, the mailbox always preserves the order in which messages arrive, i.e., FIFO

```
loop() ->  
  receive  
    dog -> ...;  
    cat -> ...;  
    snake -> ...;  
    Animal -> ...  
  end,  
  loop().
```

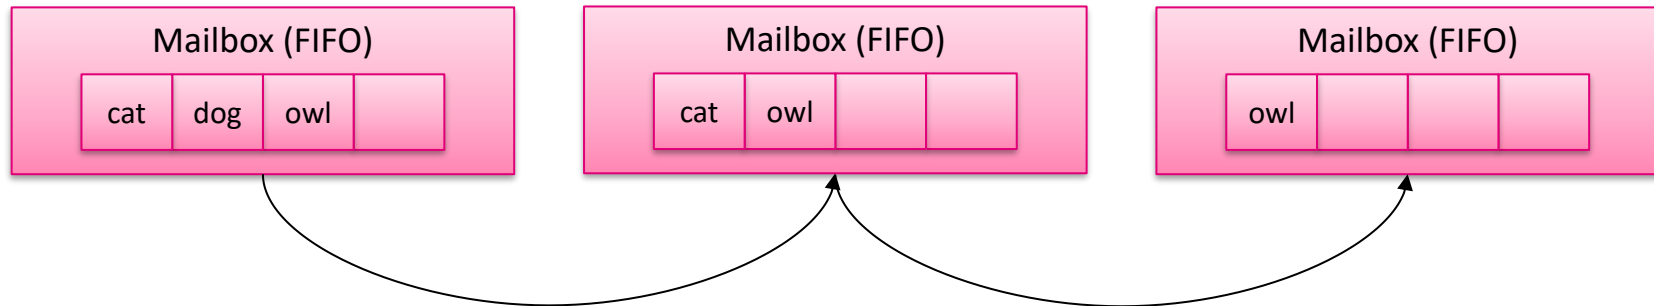


A technical note on receive in Erlang



- Messages when executing receive match from top to bottom
 - Similar to function cases
- Furthermore, the mailbox always preserves the order in which messages arrive, i.e., FIFO

```
loop() ->  
  receive  
    dog -> ...;  
    cat -> ...;  
    snake -> ...;  
    Animal -> ...  
  end,  
  loop().
```

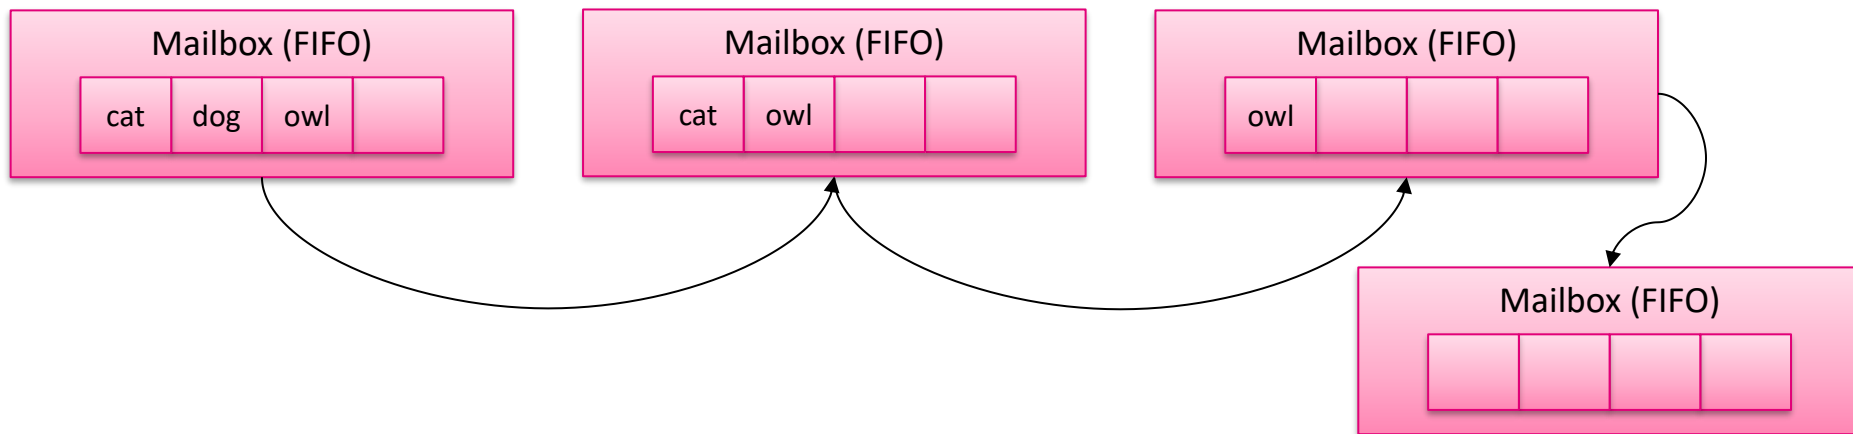


A technical note on receive in Erlang



- Messages when executing receive match from top to bottom
 - Similar to function cases
- Furthermore, the mailbox always preserves the order in which messages arrive, i.e., FIFO

```
loop() ->  
  receive  
    dog -> ...;  
    cat -> ...;  
    snake -> ...;  
    Animal -> ...  
  end,  
  loop().
```

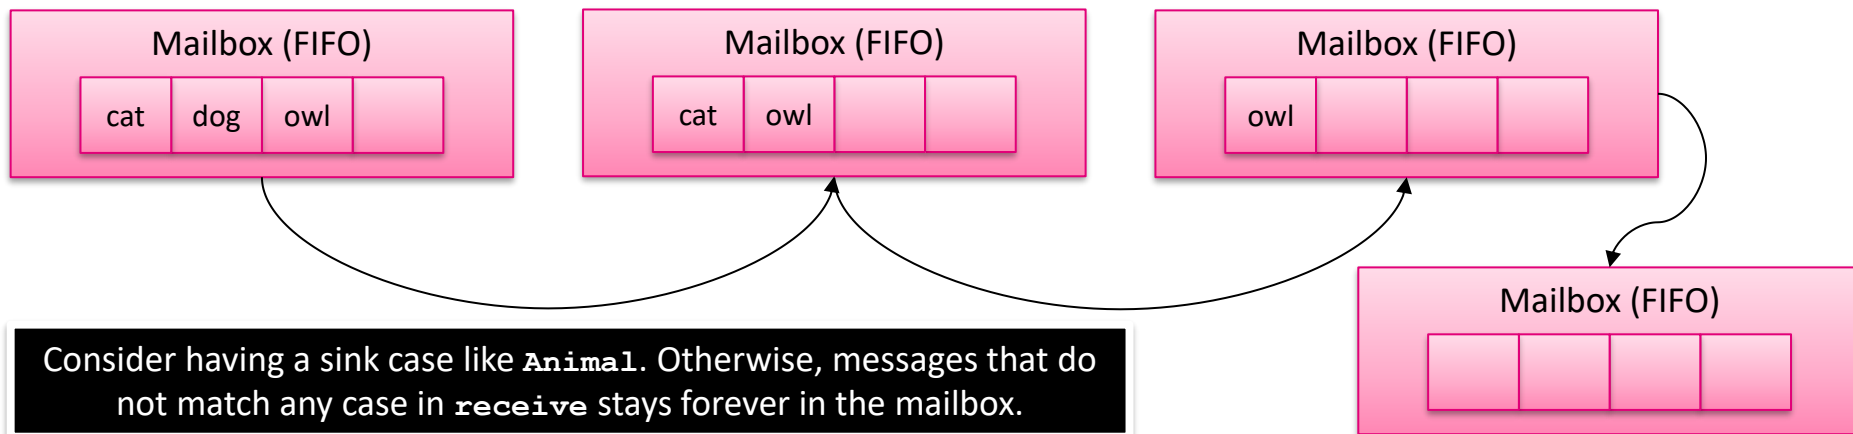


A technical note on receive in Erlang



- Messages when executing receive match from top to bottom
 - Similar to function cases
- Furthermore, the mailbox always preserves the order in which messages arrive, i.e., FIFO

```
loop() ->  
  receive  
    dog ->    ...;  
    cat ->    ...;  
    snake ->   ...;  
    Animal -> ...  
  end,  
  loop().
```



Erlang actor module – Summary



- In summary an Erlang actor module should have (at least) these elements:
 1. State
 2. Start function(s)
 3. Init function(s)
 4. Loop function
- Optionally it is also recommended to have (coming in the next slides):
 5. Message handling functions
 6. Actor API
- You may notice that the files in the code-lecture folder have the structure on the right to make it easier to write actor modules

```
-module(<module_atom>).  
-export([...]).  
  
% 1. State of the actor  
...  
  
% 2. Function(s) to create turnstile actors  
...  
  
% 3. Function(s) to initialize the state and the actors  
% behavior upon receiving messages  
...  
  
% 4. Function defining the behavior upon receiving messages  
...  
  
% 5. Message handlers  
...  
  
% 6. API  
...
```

- There is a one-to-one correspondence of the basic actor operations and concepts in Erlang

Actors Model	Erlang
Actor	Module
Mailbox Address	Process identifier (PID)
Message	Erlang term (typically an atom or tuple)
State	Erlang term (typically a record)
Behaviour	loop()
Create actor	spawn
Send message	PID ! Message
Receive message	receive ... end

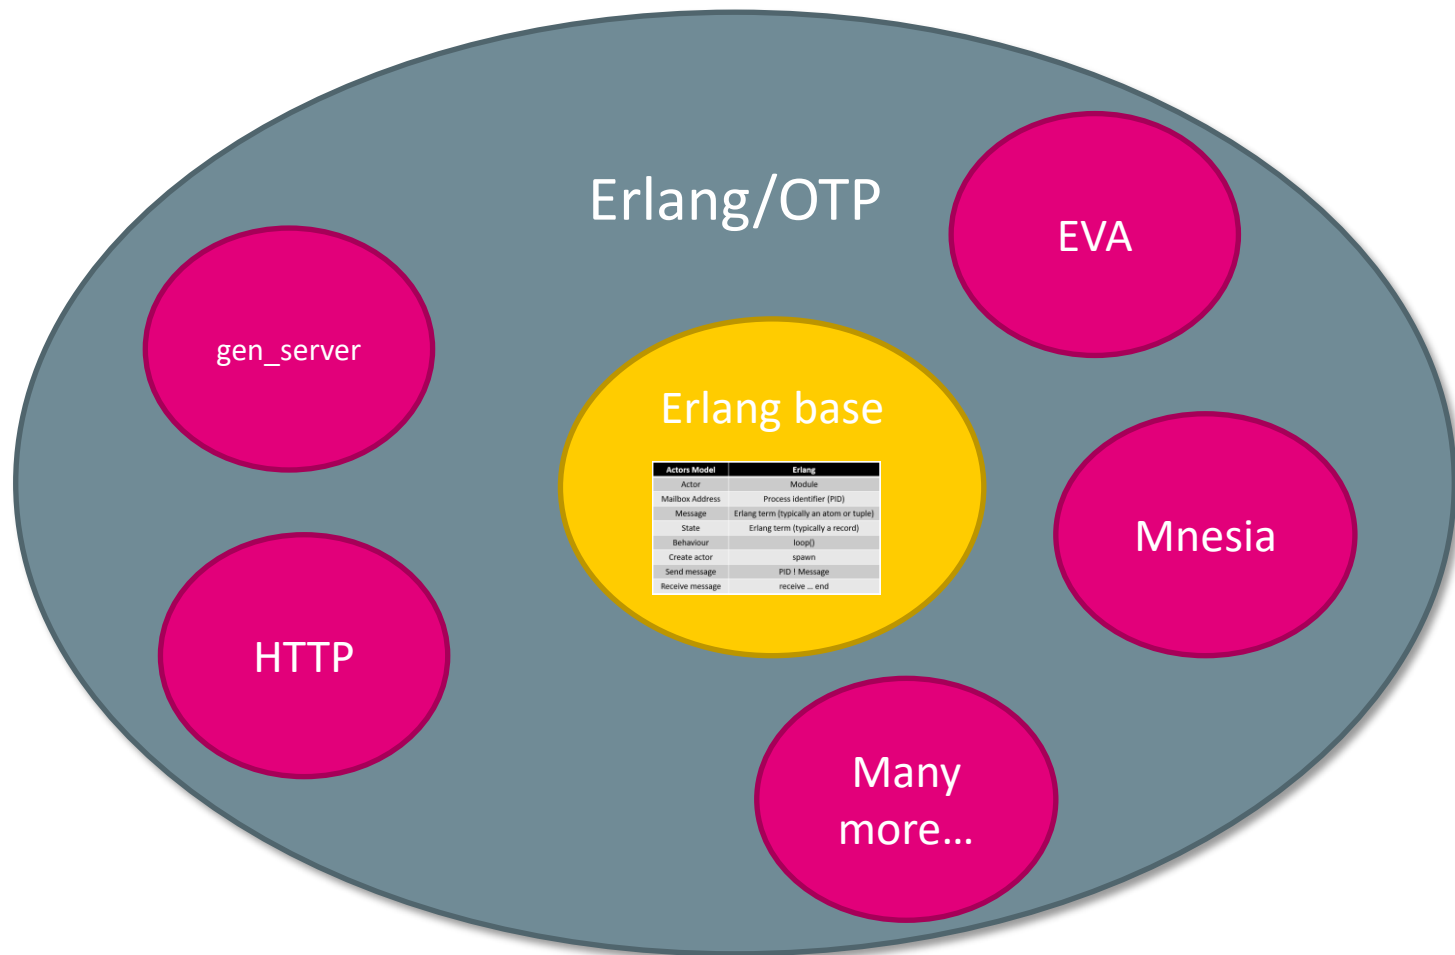
How do we use Erlang actors?



- You may interact with actors directly from the shell
 - For instance, `TurnstilePID ! person_crossing`
 - Let's try with the turnstile example
- To perform a predefined set of operations, you may write a function in a system utilities module
 - The functions in this module can be executed directly in the shell
 - This module does not necessarily need to be an actor
 - Let's look at the function `example_execution(N)` in the module `system_utilities` in folder `turnstile` for an example
- Recall that we have a guide to use Erlang for exercises
 - <https://github.itu.dk/jst/PCPP2024-Public/blob/main/general-info/guide-using-erlang-for-exercises.md>

We only used a tiny bit of Erlang

· 46

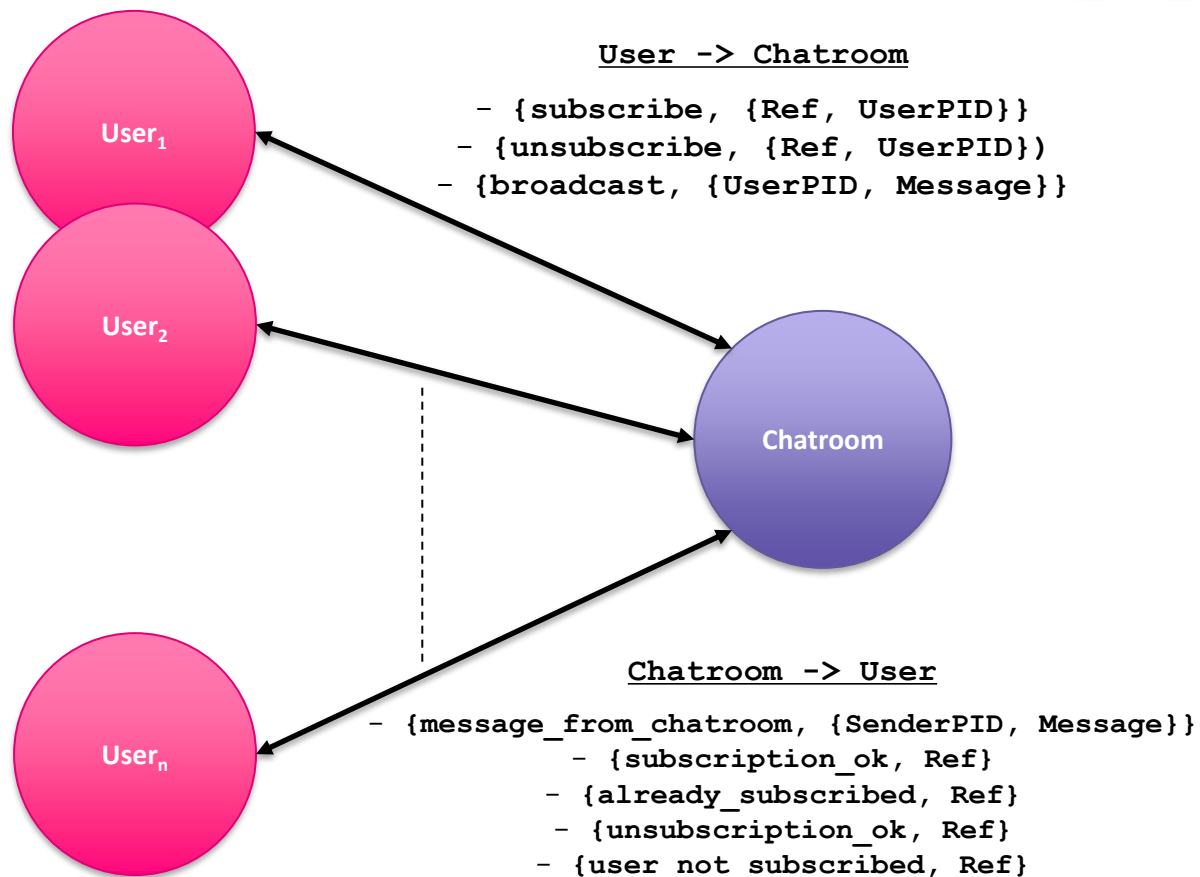




A broadcast chatroom



- A set of user actors may subscribe to a chatroom actor
 - The chatroom must confirm the subscription
- Users may emit messages that the chatroom broadcasts to all subscribers (except for the sender)
- Users may unsubscribe
 - The chatroom must confirm the unsubscription.



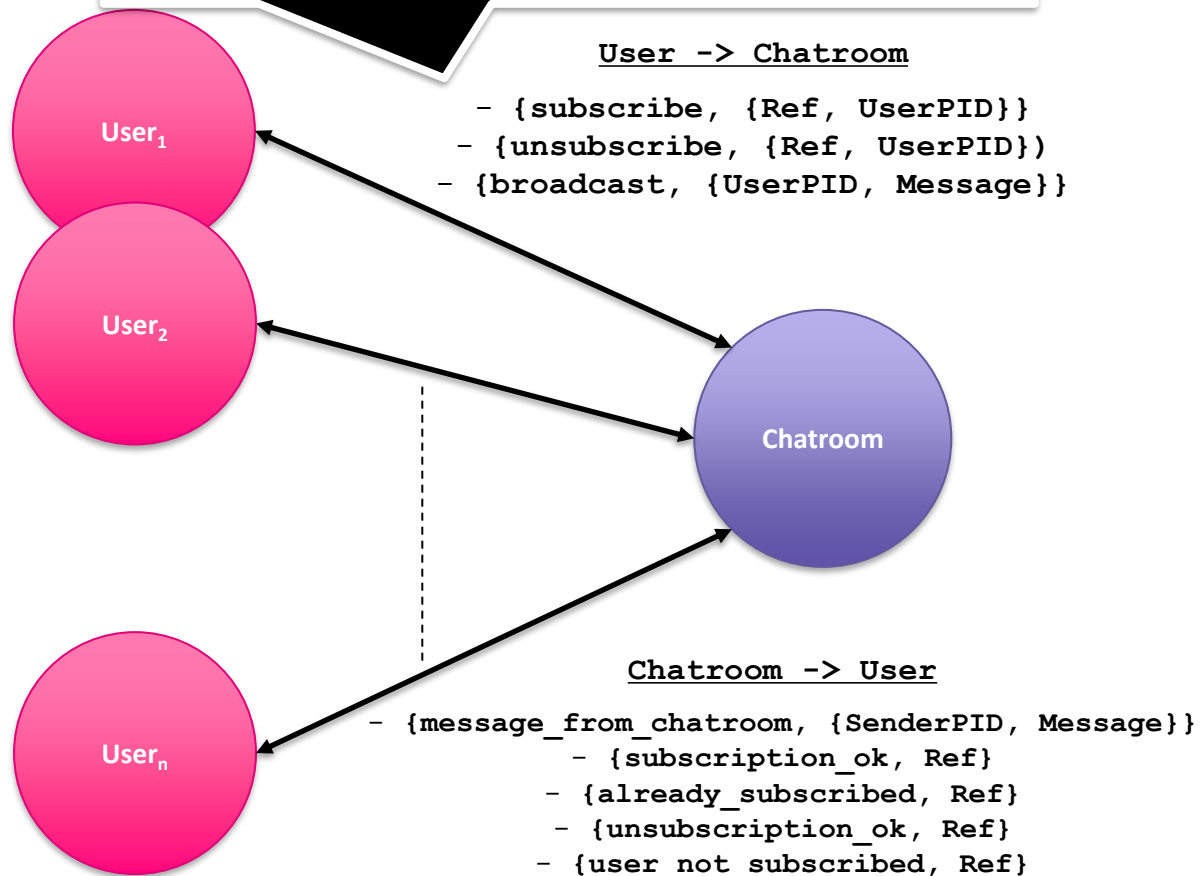
Broadcaster

· 48



Important detail, messages do not contain information about the sender. If, for instance, the sender needs a reply, the message must contain a reference to the sender

- A set of user actors may subscribe to a chatroom actor
 - The chatroom must confirm the subscription
- Users may emit messages that the chatroom broadcasts to all subscribers (except for the sender)
- Users may unsubscribe
 - The chatroom must confirm the unsubscription.



Synchronous communication - Broadcaster

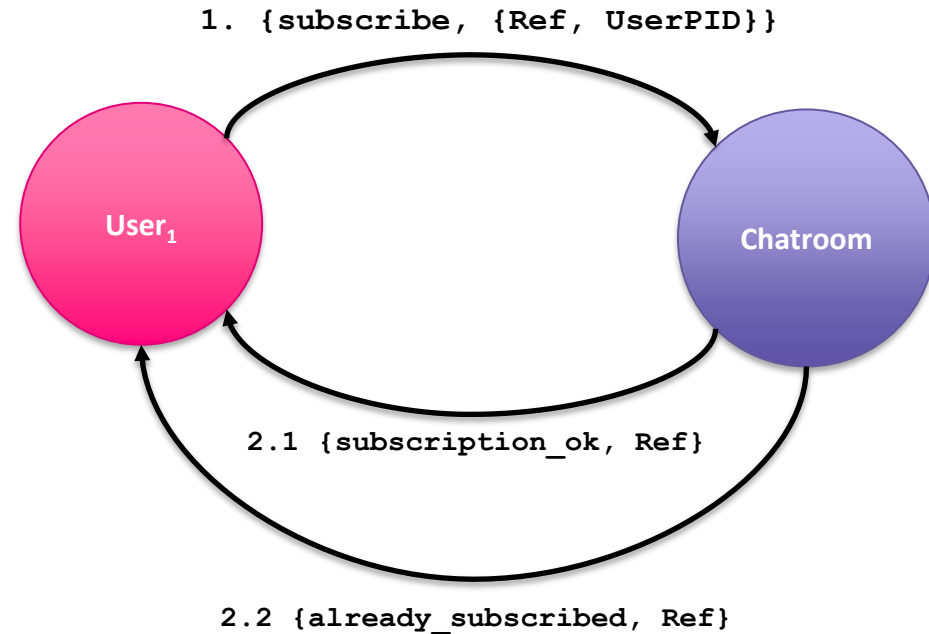


· 49

- Subscriptions require synchronous communication
 - After the user actor sends subscribe it must wait until it receives a reply from the server

```
%% User actor
handle_subscribe(ChatroomPID, State) ->
    Ref = make_ref(),
    ChatroomPID ! {subscribe, {Ref, self()}},
    receive
        {subscription_ok, Ref} ->
            NewState = State#user_state{chatroom=ChatroomPID},
            io:format("Successfully subscribed in the chatroom!~n"),
            loop(NewState);
        {already_subscribed, Ref} ->
            io:format("The chatroom says I am already subscribed.~n"),
            loop(State)
    end.

%% Chatroom actor
handle_subscribe(UserPID, Ref, State) ->
    Users = State#cr_state.users,
    case lists:member(UserPID, Users) of
        false ->
            NewState = State#cr_state{users= [UserPID|Users]},
            UserPID ! {subscription_ok, Ref},
            loop(NewState);
        true ->
            UserPID ! {already_subscribed, Ref},
            loop(State)
    end.
```



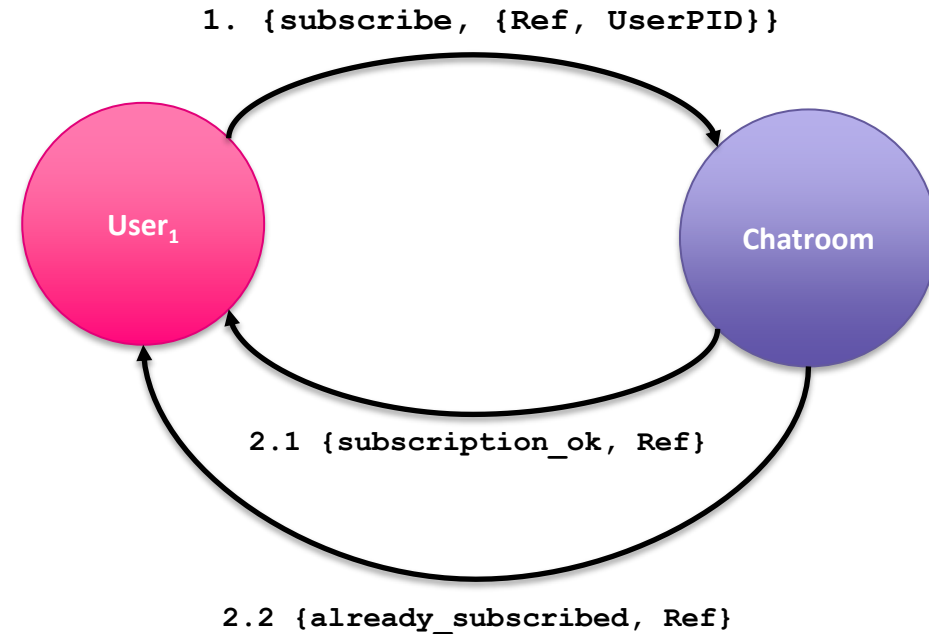
This is not a problem, as we can implement synchronous communication using the asynchronous send and receive from Erlang



- Subscriptions require synchronous communication
 - After the user actor sends subscribe it must wait until it receives a reply from the server

```
%% User actor
handle_subscribe(ChatroomPID, State) ->
  Ref = make_ref(),
  ChatroomPID ! {subscribe, {Ref, self()}},
  receive
    {subscription_ok, Ref} ->
      NewState = State#user_state{chatroom=ChatroomPID},
      io:format("Successfully subscribed in the chatroom!~n"),
      loop(NewState);
    {already_subscribed, Ref} ->
      io:format("The chatroom says I am already subscribed.~n"),
      loop(State)
  end.

%% Chatroom actor
handle_subscribe(UserPID, Ref, State) ->
  Users = State#cr_state.users,
  case lists:member(UserPID, Users) of
    false ->
      NewState = State#cr_state{users= [UserPID|Users]},
      UserPID ! {subscription_ok, Ref},
      loop(NewState);
    true ->
      UserPID ! {already_subscribed, Ref},
      loop(State)
  end.
```



This is not a problem, as we can implement synchronous communication using the asynchronous send and receive from Erlang



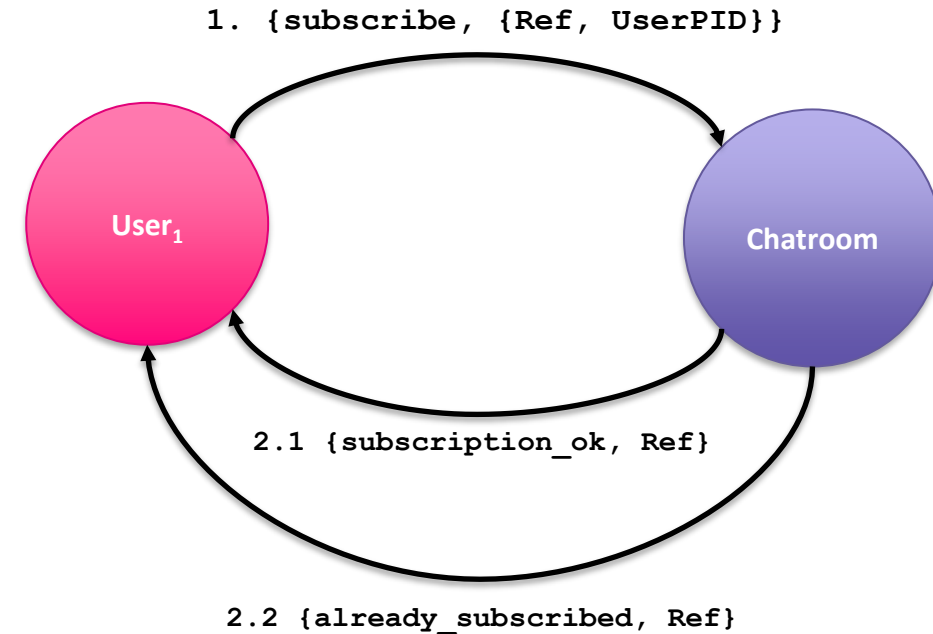
- Subscriptions require synchronous communication

References (via `make_ref()`) create a unique number that we can use to make sure that we wait for the reply to the message we sent

```
%% User actor
handle_subscribe(ChatroomPID, State) ->
  Ref = make_ref(),
  ChatroomPID ! {subscribe, {Ref, self()}},
  receive
    {subscription_ok, Ref} ->
      NewState = State#user_state{chatroom=ChatroomPID},
      io:format("Successfully subscribed in the chatroom!~n"),
      loop(NewState);
    {already_subscribed, Ref} ->
      io:format("The chatroom says I am already subscribed.~n"),
      loop(State)
  end.

%% Chatroom actor
handle_subscribe(UserPID, Ref, State) ->
  Users = State#cr_state.users,
  case lists:member(UserPID, Users) of
    false ->
      NewState = State#cr_state{users= [UserPID|Users]},
      UserPID ! {subscription_ok, Ref},
      loop(NewState);
    true ->
      UserPID ! {already_subscribed, Ref},
      loop(State)
  end.
```

When a user subscribes it must wait until it receives a



Synchronous communication

This is not a problem, as we can implement synchronous communication using the asynchronous send and receive from Erlang



· 49

- Subscriptions require synchronous communication

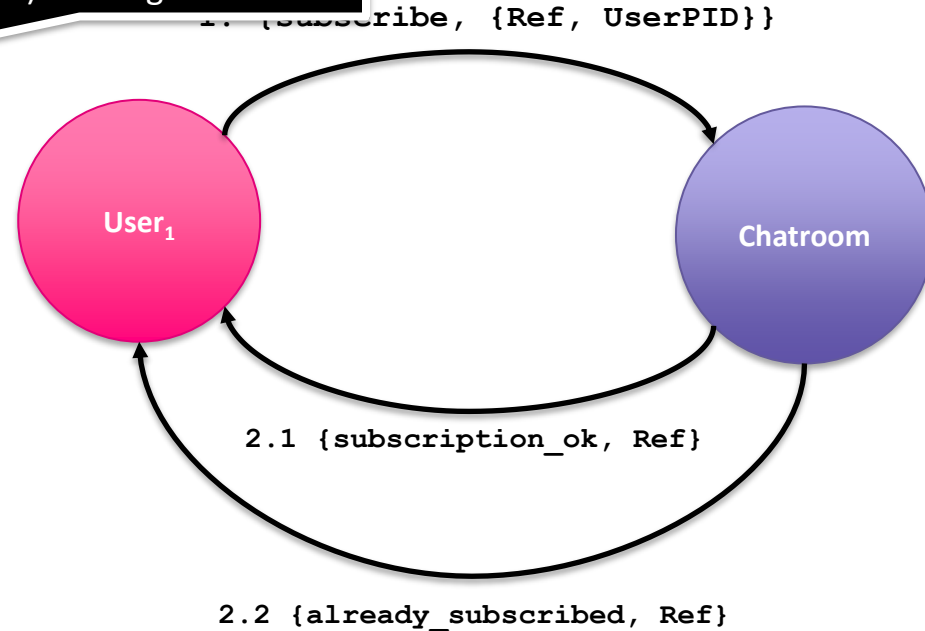
References (via `make_ref()`) create a unique number that we can use to make sure that we wait for the reply to the message we sent

The function `self()` returns the PID of the Erlang process (actor) running the code

if it receives a

```
%% User actor
handle_subscribe(ChatroomPID, State) ->
  Ref = make_ref(),
  ChatroomPID ! {subscribe, {Ref, self()}},
  receive
    {subscription_ok, Ref} ->
      NewState = State#user_state{chatroom=ChatroomPID},
      io:format("Successfully subscribed in the chatroom!~n"),
      loop(NewState);
    {already_subscribed, Ref} ->
      io:format("The chatroom says I am already subscribed.~n"),
      loop(State)
  end.

%% Chatroom actor
handle_subscribe(UserPID, Ref, State) ->
  Users = State#cr_state.users,
  case lists:member(UserPID, Users) of
    false ->
      NewState = State#cr_state{users= [UserPID|Users]},
      UserPID ! {subscription_ok, Ref},
      loop(NewState);
    true ->
      UserPID ! {already_subscribed, Ref},
      loop(State)
  end.
```



Synchronous communication - Broadcaster

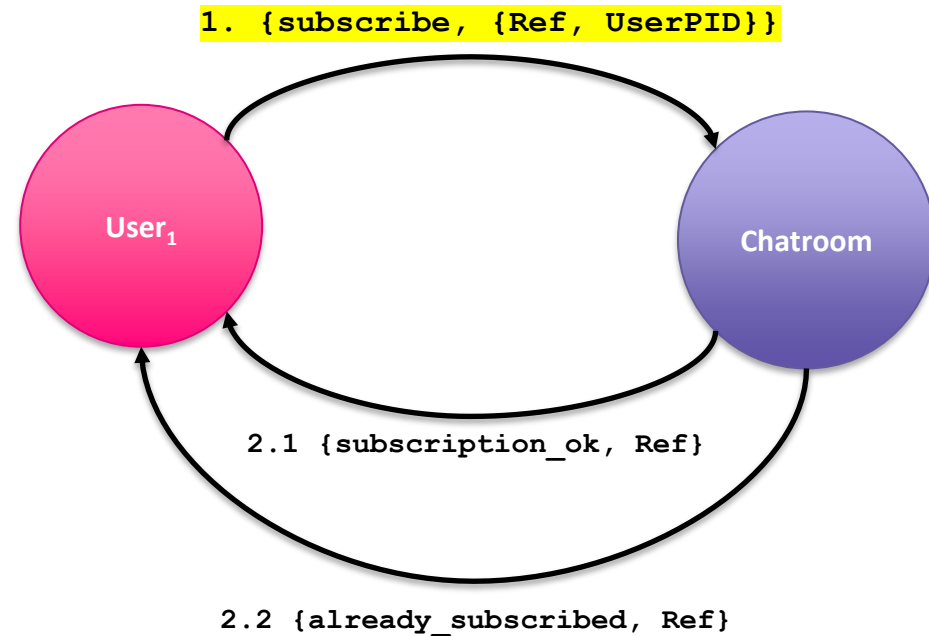


· 50

- Subscriptions require synchronous communication
 - After the user actor sends subscribe it must wait until it receives a reply from the server

```
%% User actor
handle_subscribe(ChatroomPID, State) ->
    Ref = make_ref(),
    ChatroomPID ! {subscribe, {Ref, self()}},
    receive
        {subscription_ok, Ref} ->
            NewState = State#user_state{chatroom=ChatroomPID},
            io:format("Successfully subscribed in the chatroom!~n"),
            loop(NewState);
        {already_subscribed, Ref} ->
            io:format("The chatroom says I am already subscribed.~n"),
            loop(State)
    end.

%% Chatroom actor
handle_subscribe(UserPID, Ref, State) ->
    Users = State#cr_state.users,
    case lists:member(UserPID, Users) of
        false ->
            NewState = State#cr_state{users= [UserPID|Users]},
            UserPID ! {subscription_ok, Ref},
            loop(NewState);
        true ->
            UserPID ! {already_subscribed, Ref},
            loop(State)
    end.
```



Synchronous communication - Broadcaster

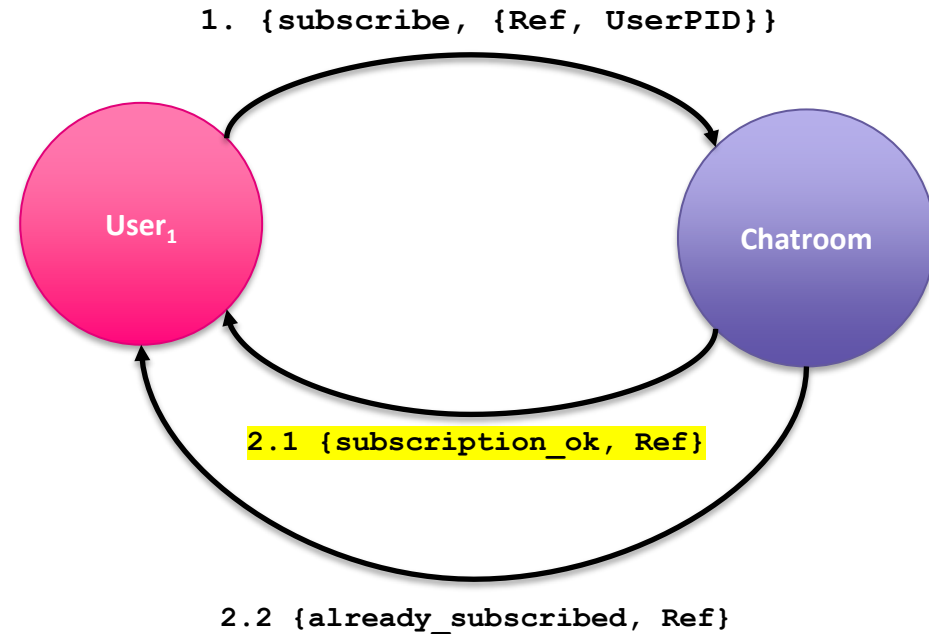


· 51

- Subscriptions require synchronous communication
 - After the user actor sends subscribe it must wait until it receives a reply from the server

```
%% User actor
handle_subscribe(ChatroomPID, State) ->
    Ref = make_ref(),
    ChatroomPID ! {subscribe, {Ref, self()}},
    receive
        {subscription_ok, Ref} ->
            NewState = State#user_state{chatroom=ChatroomPID},
            io:format("Successfully subscribed in the chatroom!~n"),
            loop(NewState);
        {already_subscribed, Ref} ->
            io:format("The chatroom says I am already subscribed.~n"),
            loop(State)
    end.

%% Chatroom actor
handle_subscribe(UserPID, Ref, State) ->
    Users = State#cr_state.users,
    case lists:member(UserPID, Users) of
        false ->
            NewState = State#cr_state{users= [UserPID|Users]},
            UserPID ! {subscription_ok, Ref},
            loop(NewState);
        true ->
            UserPID ! {already_subscribed, Ref},
            loop(State)
    end.
```



Synchronous communication - Broadcaster

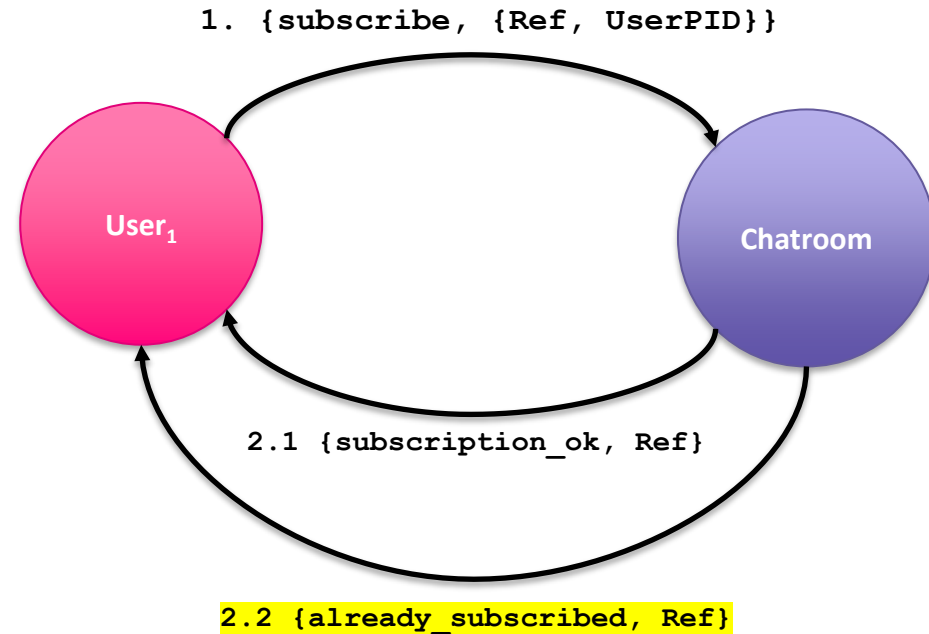


· 52

- Subscriptions require synchronous communication
 - After the user actor sends subscribe it must wait until it receives a reply from the server

```
%% User actor
handle_subscribe(ChatroomPID, State) ->
    Ref = make_ref(),
    ChatroomPID ! {subscribe, {Ref, self()}},
    receive
        {subscription_ok, Ref} ->
            NewState = State#user_state{chatroom=ChatroomPID},
            io:format("Successfully subscribed in the chatroom!~n"),
            loop(NewState);
        {already_subscribed, Ref} ->
            io:format("The chatroom says I am already subscribed.~n"),
            loop(State)
    end.

%% Chatroom actor
handle_subscribe(UserPID, Ref, State) ->
    Users = State#cr_state.users,
    case lists:member(UserPID, Users) of
        false ->
            NewState = State#cr_state{users= [UserPID|Users]},
            UserPID ! {subscription_ok, Ref},
            loop(NewState);
        true ->
            UserPID ! {already_subscribed, Ref},
            loop(State)
    end.
```



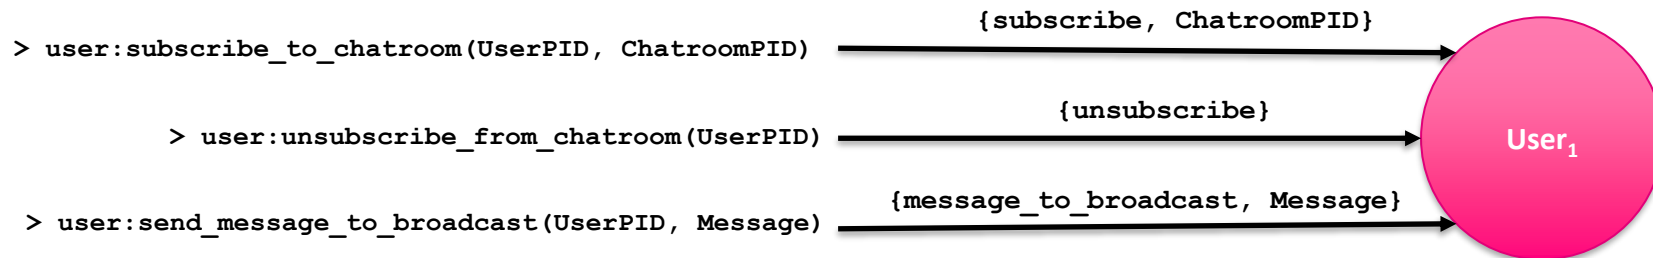
- Often it is a good idea to hide the message protocol from the users of our Erlang programs
 - To this end, define set of API functions which trigger the necessary messages

```
% 6. API
subscribe_to_chatroom(UserPID, ChatroomPID) ->
    UserPID ! {subscribe, ChatroomPID}.

unsubscribe_from_chatroom(UserPID) ->
    UserPID ! {unsubscribe}.

send_message_to_broadcast(UserPID, Message) ->
    UserPID ! {message_to_broadcast, Message}.
```

```
loop(State) ->
    receive
        {subscribe, ChatroomPID} ->
            handle_subscribe(ChatroomPID, State);
        {unsubscribe} ->
            handle_unsubscribe(State);
        {message_to_broadcast, Message} ->
            handle_message_to_broadcast(Message, State);
        ...
    end.
```



Registering processes - Broadcaster



- In Erlang, it is possible to register processes using an atom
 - This allows to send messages using said atom
 - For the broadcaster, you can use the `start_reg` function

```
%% User actor
start_reg(Username) ->
    PID = spawn(?MODULE, init, [Username]),
    register(Username, PID).

%% Chatroom actor
start_reg(ChatroomName) ->
    PID = spawn(?MODULE, init, [ChatroomName]),
    register(ChatroomName, PID).
```



Registering processes - Broadcaster




- In Erlang, it is possible to register processes using an atom
 - This allows to send messages using said atom
 - For the broadcaster, you can use the `start_reg` function

```
%% User actor
start_reg(Username) ->
    PID = spawn(?MODULE, init, [Username]),
    register(Username, PID).

%% Chatroom actor
start_reg(ChatroomName) ->
    PID = spawn(?MODULE, init, [ChatroomName]),
    register(ChatroomName, PID).
```

Demo time!




Broadcaster interesting executions



Recall: FIFO mailboxes

- Consider this execution
 1. user1 sends subscription to Chatroom
 2. user2 sends subscription to Chatroom
 3. ...

What actor will receive first subscription_ok?




Broadcaster interesting executions



Recall: FIFO mailboxes

- Consider this execution
 1. user1 sends subscription to chatroom
 2. chatroom replies subscription_ok to user1
 3. user1 emits message to Chatroom
 4. user2 sends subscription to chatroom
 5. ...

Can user2 receive the message sent by user1 in step 3?




Happened-before in distributed systems

· 59



- Note that in the previous questions the behaviour of the systems depends on the reception of messages
- Thus, the happened-before relation defined by Lamport is useful in reasoning about actor systems
 - An action a happens-before an action b if they belong to the same actor and a was executed before b
 - A $\text{send}(m)$ action happens-before its corresponding $\text{receive}(m)$
- Note the similarity with the happens-before relation of the Java memory model
 - We reason about message exchange instead of locking (but *inherent coordination problems remain*)
 - Visibility issues disappear as actors only access local memory



A bounded buffer

Producer-consumer problem | Intuition



· 61

- Perhaps more intuitive example

Consumers

Producers



Shared data structure of fixed size

Producer-consumer problem | Intuition

· 62



Consumers

Producers



Shared data structure of fixed size

Bounded Buffer with Actors

· 63



Consumer -> Buffer

- {get, {Ref, SenderPID}}

Producer -> Buffer

- {put, {Ref, Elem, SenderPID}}

Bounded
Buffer

Buffer -> Consumer

- {get_ok, {Ref, elem}}

Consumer_i

Buffer -> Producer

- {put_ok, Ref}

Producer_i

Bounded Buffer with Actors

· 63



Consumer -> Buffer

- {get, {Ref, SenderPID}}

Producer -> Buffer


- {put, {Ref, Elem, SenderPID}}

Buffer -> Consumer

- {get_ok, {Ref, elem}}

Buffer -> Producer

- {put_ok, Ref}



Let's look at the code
(bounded_buffer package)

- After a request to consume or produce, the actors must wait for the reply from the bounded buffer (i.e., synchronous communication)
- Consumers wait if the buffer is empty
- Producers wait if the buffer is full


Deadlocks in Actor Systems – Bounded Buffer



· 64

- An actor system may end up in a deadlock state

```
deadlock_situation() ->
  BufferPID = buffer:start(1),
  A1 =
    fun () ->
      buffer:get(BufferPID),
      buffer:put(BufferPID, something),
      io:format("Done~n")
    end,
  A2 =
    fun () ->
      buffer:get(BufferPID),
      buffer:put(BufferPID, something),
      io:format("Done~n")
    end,
  lists:foreach(fun (A) -> spawn(A) end, [A1,A2]).
```



Actors in distributed systems



The actors model has natural mapping in distributed systems

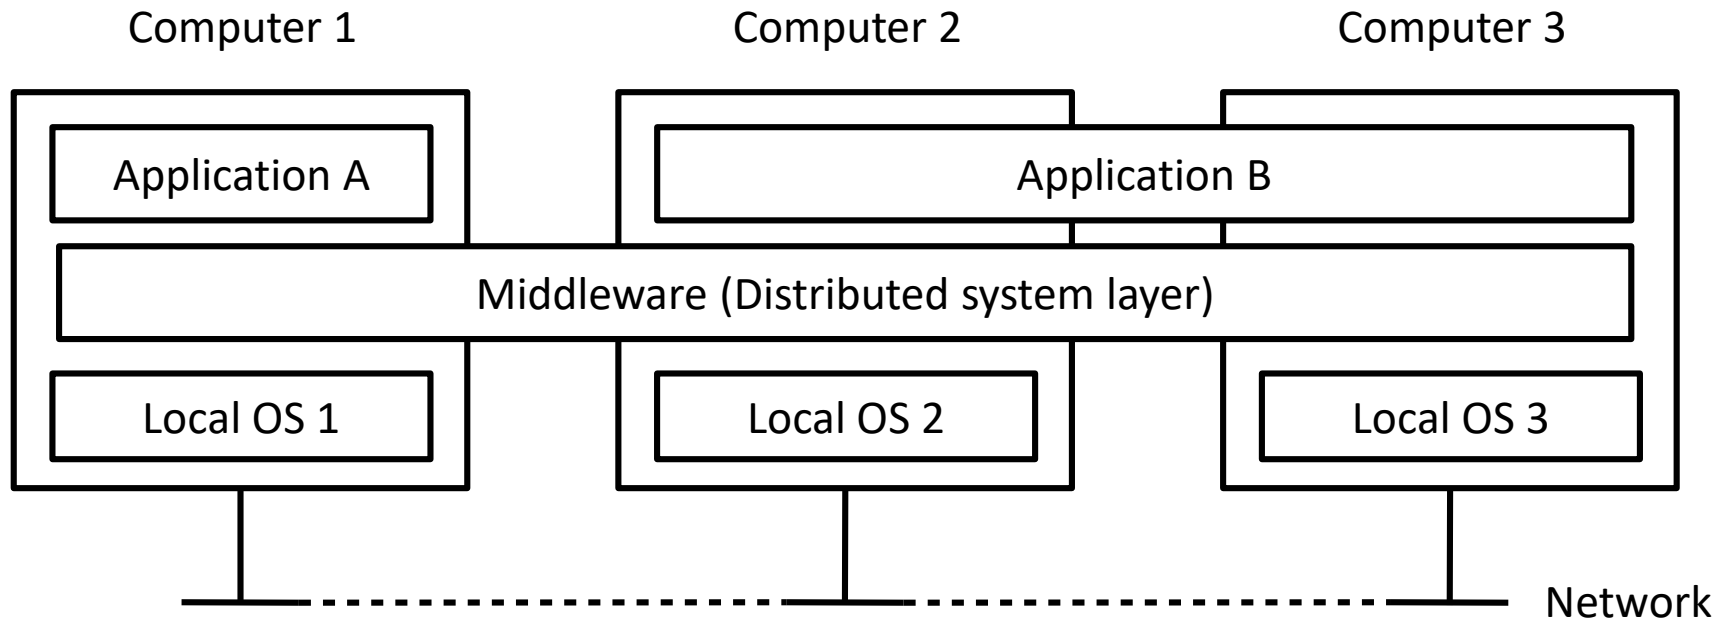
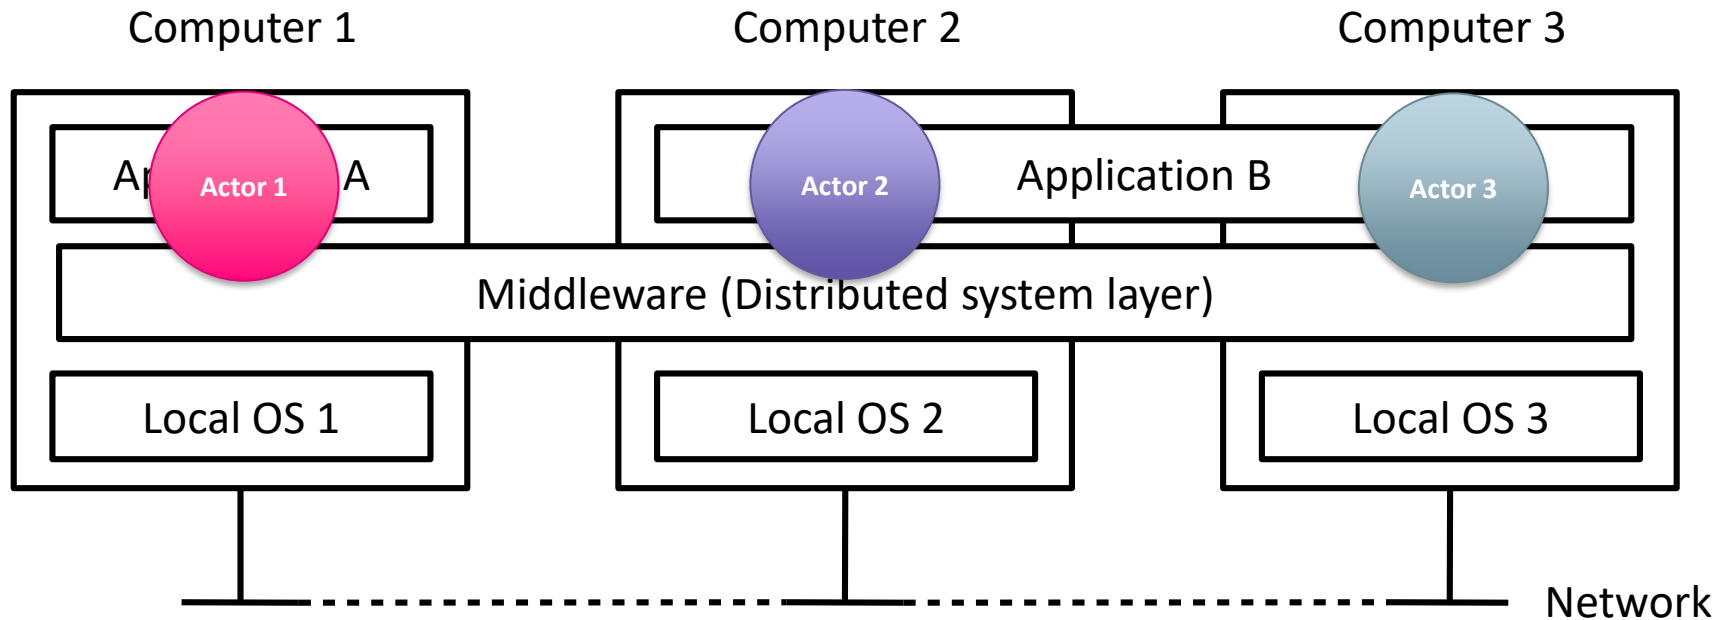


Figure taken from -> Distributed Systems: Principles and Paradigms. Andrew S. Tanenbaum and Maarten Van Steen. 2007.

Actors in distributed systems



The actors model has natural mapping in distributed systems




Actors in a single computer



· 67

The actors model is applicable in a single computer as well



In this lecture, we focus on this type of actor system

- Problems in shared memory concurrency (revisited)
- Actors
- Erlang
- Example systems
 - Turnstile (counter)
 - Broadcaster
 - Bounded Buffer